

Does a Toehold Make a Bidder Bolder?

Preemption and Multiplicity in Multi-Round Takeover Auctions

Zain Naboulsi

Principal AI Engineer, Sparq

zain.naboulsi@teamsparq.com

A *toehold*, a stake in a target quietly acquired before a takeover contest, is supposed to make a bidder bid more aggressively and scare off rivals. Theory has long predicted this advantage, yet toeholds are rare in practice [Betton et al., 2009]: the “toehold puzzle.” We ask whether the predicted deterrence survives in a contest that actually unfolds over several rounds, rather than the one or two moves earlier models assume. Building on our companion study of sealed takeover auctions [Naboulsi, 2026], we cast a multi-round ascending common-value takeover auction with a toehold as a two-player imperfect-information game and solve for its own-profit Bayes-Nash equilibria on one small instance, certified as ε -equilibria with ε at most 8×10^{-5} and typically nearer 10^{-6} , using a forced-opening best-response construction that reads the value of a preemptive jump bid without trusting any single solver run. Three findings reshape the picture. First, the game has *multiple* equilibria that give the toehold-holder the same profit to three decimals, and to four at two of the three toeholds we certify, yet differ completely in conduct: in one the rival reads a jump bid as strength and folds a third of the time, in another it folds far less often, or at a zero toehold essentially never, and no one preempts. So deterrence is not pinned down by the toehold. Second, and decisively, preemption occurs even at a *zero* toehold, so preemption is not something a toehold buys. Third, the clean “bigger toehold, more deterrence” relationship that classical two-move models predict is real at exactly one round and *dissolves* as the contest lengthens: raising the toehold from 15% to 30% lifts deterrence sharply, by about three-quarters, in a single-round, two-move contest, buys nothing once the contest has a genuine second round, and by the third round the equilibria multiply and deterrence is no longer pinned down at all. In our model the toehold’s deterrence channel is an artifact of the two-move approximation, which speaks to the toehold puzzle: the bidding edge that would justify taking a toehold is not identified once the contest has more than one round. On the computational side we benchmark deep and exact solvers on the sequential game, finding that a generic policy-gradient method (PPG) reaches a near-exact equilibrium while a standard imperfect-information baseline (NFSP) fails to converge on the same budget, and we extend the analysis into a multi-round multi-signal regime too large to enumerate, where exact solvers cannot run. We release the games, solvers, and experiments.

 <https://github.com/zainnab-sparq/sequential-takeover-auctions>

Sparq

Keywords: sequential takeover auctions, toeholds, preemptive bidding, equilibrium multiplicity, imperfect-information games, policy gradient, exploitability, OpenSpiel

1 Introduction

A bidder who already owns a slice of a target enters a takeover contest differently from one who owns nothing. That slice, a *toehold*, changes the arithmetic of winning and losing: the toehold-holder collects a share of whatever price the eventual winner pays, so it is willing to bid more aggressively, and a rival, anticipating that aggression, may decline to compete at all. This is the toehold’s celebrated deterrence or “aggressiveness” channel [Bulow et al., 1999, Fishman, 1988]. It is also the source of a long-standing empirical puzzle: if a toehold confers a bidding edge, bidders should acquire toeholds routinely before launching a bid, yet Betton et al. [2009] document that they seldom do. Explanations invoke the cost of tipping off the market, legal disclosure thresholds, or the risk of being seen as hostile; Dai et al. [2021] argue instead that toeholds are effective once the selection into taking one is corrected for, so the puzzle is partly an endogeneity artifact. We ask a different question: is the predicted bidding edge even there, once the contest is modeled as it actually plays out, over several rounds of escalation rather than a single sealed bid or a two-move exchange?

Our companion paper [Naboulsi, 2026] studied *sealed*, simultaneous takeover auctions, where each bidder submits one bid and the toehold’s aggressiveness channel is competed away, so a sealed model cannot express preemption at all. That paper asserted, but could not show, that the channel “lives in sequential, asymmetric contests.” This paper makes the sequential contest the object of study and turns the assertion into a result, though not the result we expected. We model a multi-round ascending common-value auction in which a toehold-holder and a rival alternate, each able to jump to any higher price on a discrete grid or to pass, with the bid history public and each bidder’s due-diligence signals private. We solve for the game’s own-profit Bayes-Nash equilibria on small instances, to certified ε -equilibria, and, crucially, measure the economic value of a preemptive jump bid with a construction that does not depend on any solver converging to a particular equilibrium.

The results overturn the simple story in three steps. Equilibria are *multiple*: independent solves land on the same profit for the toehold-holder, to four decimals at two of the three toeholds we certify and to within 6×10^{-5} at the third, and on completely different conduct, one preempting and one not, because “I preempt because you fold, you fold because I preempt” and “I open low because you will not fold” are both self-consistent. Preemption survives at a *zero* toehold, so it is a feature of the sequential structure, not a purchase the toehold makes. And the tidy monotone relationship between toehold size and deterrence that classical two-move models deliver is an artifact of stopping at two moves: it holds at one round, saturates at two, and by three rounds the equilibria multiply and deterrence is no longer pinned down at any toehold. The channel does not decay gradually as the contest lengthens; it is gone by the second round.

We make three contributions. (1) To our knowledge this is the first computational solve, to certified ε -equilibria, of a genuinely multi-round ascending common-value takeover auction with a toehold, including a forced-opening best-response tool (itself exact) that reads preemption incentives without trusting solver convergence. (2) We establish equilibrium multiplicity and the two-move artifact as the central economic findings, reconciling them with the signalling account of preemptive bidding [Fishman, 1988] and with the toehold puzzle. (3) On the computational side we benchmark deep and exact solvers on the sequential game and extend to a multi-round, multi-signal regime beyond exact enumeration, showing which methods still work when the game outruns the exact solvers. All games, solvers, and experiments are released.

2 Related work

Toeholds and preemptive bidding. The idea that a toehold makes a bidder more aggressive, lowering a rival’s chance of winning, is classic [Bulow et al., 1999, Klemperer, 1999].

Preemptive *jump* bidding, bidding above the minimum required to signal strength and deter entry, was formalized by Fishman [1988] as a *signalling* phenomenon: a high bid credibly conveys a high valuation, and a rival holding a weak signal infers that competing means overpaying, so it withdraws. Crucially, Fishman’s mechanism requires no toehold. Avery [1998] showed jump bidding can arise even in a setting with no private information to signal, as a coordination device. Daniel and Hirshleifer [2018] develop a costly sequential bidding model in which each bid entails a cost, generating preemptive first bids and infrequent competition. Closest to our setting, Dodonova [2012] studies toeholds and signalling in takeover auctions in a compact model; our multi-round treatment reproduces the classical comparative static, that a larger toehold deters more, in a one-round contest, and then shows it does not survive additional rounds. Empirically, toeholds are rarer than the classical edge would predict (the toehold puzzle), a tension our multiplicity and two-move-artifact findings speak to directly.

Common-value auctions and the winner’s curse. Bidding in a common-value auction requires shading against the adverse selection of winning [Capen et al., 1971, Milgrom and Weber, 1982, Wilson, 1977]. Our game is common-value with noisy private signals, and the preemption mechanism we find is a winner’s-curse inference: a rival that reads a jump bid as evidence of a high common value folds, because competing would mean winning mainly in the states where its own signal is misleadingly high. This is the same inference as in Fishman [1988], realized computationally.

Computational game solving. Exact and sampling solvers for imperfect-information games, counterfactual regret minimization [Zinkevich et al., 2007, Brown et al., 2019], magnetic mirror descent [Sokota et al., 2023], and policy-space response oracles [Lanctot et al., 2017], have mostly been developed on recreational benchmarks. Generic deep reinforcement learning, PPO [Schulman et al., 2017], PPG [Cobbe et al., 2021], and NFSP [Heinrich and Silver, 2016], has been argued to rival specialized machinery on such games [Rudolph et al., 2026], and multi-agent reinforcement learning is increasingly used to analyze auction mechanisms directly [d’Eon et al., 2024]. Sequential, asymmetric games are the setting where specialized regret-based solvers are expected to have the largest edge, so they are a sharp test. We benchmark both families on our sequential auction and, following our companion paper, use a learned-best-response exploitability estimate [Lanctot et al., 2019] when the game grows beyond exact enumeration.

3 A multi-round takeover auction with a toehold

Primitives. Two risk-neutral bidders compete for a target of common value W , drawn uniformly from a grid of `num_values` levels. Each bidder i receives $k = \text{num_signals}$ noisy private signals of W ; with probability $1 - \text{signal_noise}$ a signal equals W and otherwise is uniform on the value grid. Bidder 0 holds a toehold $\theta \in [0, 1)$: a fraction of the target it already owns. Prices lie on a discrete grid of `num_bids` levels spaced by `bid_step`.

Protocol. The auction is an ascending, alternating-move contest of up to $R = \text{num_rounds}$ rounds. A round is one turn for each bidder, so an R -round contest is $2R$ moves and $R = 1$ is the two-move exchange, one bid and one response, that the classical models analyze. We count in rounds throughout. On its turn a bidder either *raises*, jumping to any grid price strictly above the standing bid (a jump bid is thus available, not just the minimum increment), or *passes*. Passing while behind ends the auction and the standing bidder wins; both passing with no standing bid ends it with no sale; otherwise play continues until the round limit. The bid history (the standing bid, whose turn it is, and the round index) is public and carried in each bidder’s information state; a bidder’s own signals stay private; the game has perfect recall. Because raises must strictly increase, the number of raises is bounded by the price grid: the

tree depth is capped by `num_bids`, not by R , so added rounds deepen the contest only while `num_bids` $\gtrsim 2R$.

Payoffs. If bidder 1 wins at price p it earns $W - p$ and the toehold-holder earns θp (it sells its share into the winning price). If the toehold-holder wins at price p it earns $W - (1 - \theta)p$, paying only for the fraction it does not already own. The setup matches the payoff contract of our companion sealed game, so the two are directly comparable on an identical grid. Note that $R = 1$ is *not* that sealed game: the rival still moves second and still sees the opening bid, which is exactly why deterrence is a meaningful quantity at $R = 1$ and is not one in a sealed auction.

The instance. All the economics below is one game: `num_values` = 3, `num_bids` = 9 with `bid_step` = 0.375 (so the grid spans $[0, 3]$, the full range of W), $k = 1$ signal per bidder, `signal_noise` = 0.5 for both bidders, and R swept over $\{1, 2, 3\}$. Each cell is solved from 6 restarts (restart 0 uniform, the rest drawn from seed 20,260,713). The certified runs use a 2×10^6 -iteration budget with the tolerance set to 10^{-8} , and a restart enters a table only if its own-profit NashConv is at or below 10^{-4} ; both the budget and that acceptance bound matter, and both are recorded per row in the released data. The benchmark, scaling and intractable-regime instances of Sections 8 and 9 are different and smaller, and are stated there.

Encoding. The game is implemented as an OpenSpiel [Lanctot et al., 2019] extensive-form game with the bid history as public observation tokens and the private signals as one-hot features. Perfect recall and the information-set structure (own signals hidden, bid history public) are enforced by construction and checked by unit tests; exact counterfactual regret minimization drives exploitability below 0.01 on small instances, confirming the solver converges on the encoding as expected.

4 Solving for own-profit equilibria

Standard solvers optimize the zero-sum rendering of a game. The economics we care about live in each bidder’s *own profit*, a general-sum object. We therefore solve for own-profit Bayes-Nash equilibria with extensive-form fictitious play: at each iteration each bidder computes an own-profit best response by traversing the tree and scoring terminals with raw profits (not zero-sum returns), then behavioral strategies are averaged and the own-profit NashConv (the summed best-response gain) is tracked as an ε -equilibrium certificate. General-sum Nash is PPAD-hard, so no solver is owed a unique answer; we treat convergence as an empirical property to be verified, not assumed.

Reading economics without trusting the solver. A comparative static read off a single fictitious-play solve is unreliable, because the solve silently selects one equilibrium out of possibly many. We therefore add a *forced-opening best-response* construction: we pin a bidder’s opening information sets to a chosen bid while letting every later decision best-respond exactly, and sweep the forced opening across the price grid to trace the profit of each possible jump bid. Because each information set lives at a single tree depth, this pins exactly the opening and nothing else. The construction is validated against a brute-force enumeration over every pure opening policy, and it measures what a preemptive jump is worth *without any reliance on fictitious play converging*, breaking the circularity of reading economics off a solve that is itself in question.

5 Equilibrium multiplicity: same value, different deterrence

The central empirical fact is that in the sequential toehold auction the equilibrium *value* is pinned but the equilibrium *behavior* is not. Independent fictitious-play restarts converge to the same own profit for the toehold-holder, to four decimals at $\theta = 0$ and $\theta = 0.15$ and to within 6×10^{-5} at $\theta = 0.05$, and to completely different conduct. Table 1 reports representative pairs. At a toehold of 0.05, two profiles worth 0.2716 and 0.2715 to the toehold-holder deter the rival with probability 0.002 and 0.333 respectively; the preempting profile is certified as an ε -Nash equilibrium with $\varepsilon \approx 4 \times 10^{-6}$ and the passive profile, which converges more slowly, with $\varepsilon \approx 8 \times 10^{-5}$.

The contrast is not uniform across the toehold range, and where it fades it does so for two different reasons. It is sharpest at $\theta = 0$ (0.000 versus 0.333) and narrows to 0.220 versus 0.333 by $\theta = 0.15$. At $\theta = 0.20$ every accepted restart lands on the same conduct (0.3333 to four decimals), so there is nothing to certify. Above $\theta = 0.20$ no restart converges in the base three-round game at the budget of the nine-point sweep, so deterrence is unidentified there by non-convergence rather than by multiplicity: four of those nine toeholds (0.25, 0.30, 0.40, 0.50) yield no usable equilibrium and are absent from the released summary.

Toehold θ	Bidder-0 value		Rival deterrence	
	profile A	profile B	profile A	profile B
0.00	0.1840	0.1840	0.000	0.333
0.05	0.2716	0.2715	0.002	0.333
0.15	0.4554	0.4554	0.220	0.333

Table 1: Equilibrium multiplicity in the three-round auction. Each row shows two profiles, each an ε -Nash equilibrium, that give the toehold-holder the same value but deter the rival with very different probability. Certified ε by row: 2.4×10^{-6} at $\theta = 0$, 6.4×10^{-6} at $\theta = 0.15$, and 7.6×10^{-5} at $\theta = 0.05$, where the passive profile is the slower of the pair (its partner is at 4.3×10^{-6}). Values agree to four decimals except at $\theta = 0.05$, where they differ by 6×10^{-5} , less than that row’s ε , so the two values are not distinguishable at the precision we certify. Deterrence is not identified by the toehold. Regenerate with `sequential_general_sum.py certify`, or in one second from the committed per-restart data with `certificate_from_selection_csv`.

Why, mechanically. The forced-opening profit curve shows both conducts are self-consistent (Figure 1). The cleanest case is the one that carries the second finding. At $\theta = 0$ the two equilibria are a rival that folds with probability 0.333 and a rival that essentially never folds (3×10^{-5}). Against the first, committing to the jump at 1.125 rather than opening at the bottom of the grid raises the toehold-holder’s profit by 179%; against the second it raises it by exactly nothing, and the curve peaks at the cheapest opening. So “I preempt because you fold, you fold because I preempt” and “I open low because you will not fold, you will not fold because I open low” are both equilibria. This is coordination-style multiplicity, not solver failure.

The same pattern appears wherever we can test it, and we can test it less often than we would like. Against the most-detering equilibrium rival the jump premium is large at every toehold, ranging from 31% at $\theta = 0.15$ to 179% at $\theta = 0$. Against the least-detering one it collapses to 2.6% and 5.3% at $\theta = 0.10$ and 0.15, where that rival folds with probability 0.17 and 0.22. But at $\theta = 0.05$ and $\theta = 0.20$ the least-detering equilibrium we find is not a passive rival at all: it still folds with probability 0.25 and 0.33, and the premium against it stays at 77% and 25%. Those two toeholds therefore do not test the contrast; only $\theta = 0$ exhibits a genuinely non-folding equilibrium rival. As a control, against a *uniform*, unbluffable rival the forced-opening curve slopes straight down with a zero jump premium at every toehold: a jump bid buys deterrence, and deterrence is worthless against a rival that cannot be deterred. Premia

here are measured against the profit of opening at the bottom of the grid; the released data reports that normalization and the peak-relative one in separate columns.




Figure 1: Forced-opening best-response profit for the toehold-holder as a function of the opening bid it commits to, with every later decision best-responding. Left: against the most-detering equilibrium rival at each toehold, profit peaks at the preemptive jump. Middle: against the least-detering equilibrium rival. Only the $\theta = 0$ curve there is a genuinely non-folding rival, and it is the one that peaks at the cheapest opening; at the other toeholds the least-detering equilibrium still folds with probability 0.17 to 0.33, and its curve still peaks at the jump, which is what the panel subtitles report. Right, the control: against a rival bidding at random, which cannot be deterred, profit slopes straight down with no jump premium at any toehold. A jump bid only pays when it can scare a rival off, which is why both the preempting and the non-preempting conduct are self-consistent equilibria. The forced-opening best response is exact; the two equilibrium rivals are solver output, and in the released version of this figure they are default-budget solves that stopped at NashConv 10^{-4} rather than probe-precision ones (see Section 10). Only the control panel is solver-free.

Preemption at a zero toehold. The decisive observation is in the first row of Table 1. At $\theta = 0$, some restarts still converge on a profile in which the toehold-holder jump-bids and the rival folds a third of the time, certified at $\varepsilon \approx 2.4 \times 10^{-6}$. *Preemption is not something the toehold buys.* The toehold sets the toehold-holder’s profit; the sequential structure is what makes preemption possible; which equilibrium is played is a matter of coordination. The original hypothesis, that a toehold buys a preemptive jump bid and revives the aggressiveness channel, is false as a causal claim.

A methodological caution. Reaching this conclusion required checking the full set of restarts, not a sample. An early analysis re-ran three restarts at a tight tolerance, saw them agree, and wrongly concluded a competing equilibrium was a numerical artifact; two other restarts converged to it at $\varepsilon \approx 6 \times 10^{-7}$. When testing for multiplicity, a sample of restarts that agree proves nothing, because the disagreeing basin may simply not be in the sample. Likewise, a fictitious-play run that stops *at* its tolerance has not converged; it has been let go. Every number in Tables 1 and 2 was re-run with the tolerance set far below what the run can reach, so the compute budget, not the stopping rule, ended it. The one exhibit that predates this discipline is the mechanism figure, whose rival policies are default-budget solves; the `certify` entry point now re-runs that study too, and we flag the gap rather than let the claim cover it.

6 The two-move artifact

Sweeping the number of rounds turns the project around. Table 2 reports deterrence as a function of the toehold at one, two, and three rounds. Every cell with a number in it is certified

at ε on the order of 10^{-6} to 10^{-5} (verified with the tolerance set to 10^{-8} and the run ended by budget, far tighter than the default stopping rule); the one cell without a number is the one where no restart reaches the acceptance bound at all, and it is reported as such.

Toehold θ	Rival deterrence probability		
	$R = 1$ round	$R = 2$ rounds	$R = 3$ rounds
0.15	0.333 (unique)	0.333 (saturated)	0.220 or 0.333 (multiple)
0.30	0.583 (unique)	0.333 (saturated)	unconverged

Table 2: The toehold-to-deterrence channel by number of rounds. At one round the equilibrium is unique (at a positive toehold) and deterrence rises with the toehold, exactly as classical two-move intuition predicts. At two rounds deterrence saturates and stops responding to the toehold. By three rounds the equilibria multiply and deterrence is no longer pinned by the toehold: it spans 0.220 to 0.333 at $\theta = 0.15$, and 0.000 to 0.333 at a zero toehold (a row not shown here), while at the largest toehold no restart converges even at the tightest budget.

At *one* round, and a positive toehold, the equilibrium is unique: independent restarts land on identical numbers, and deterrence is markedly higher at the larger toehold (0.333 at $\theta = 0.15$, 0.583 at $\theta = 0.30$, with all six restarts identical at each). We sweep only these two positive toeholds at $R = 1$, so this is an increase between two points rather than a monotone curve. This reproduces the classical comparative static, in our model, at one round. At *two* rounds deterrence *saturates*: raising the toehold from 0.15 to 0.30 moves deterrence not at all (0.333 in both cases), whereas a two-move contest bought a large jump ($0.333 \rightarrow 0.583$). By *three* rounds the equilibria multiply, deterrence spans 0.000 to 0.333 across profiles and toeholds (and 0.220 to 0.333 at $\theta = 0.15$ alone), and preemption is available even at a zero toehold; at the largest toehold ($\theta = 0.30$) no restart converges at all, even at the tightest budget, so deterrence is not identified there by any route.

So the toehold’s deterrence channel behaves exactly as auction theory says *in a two-move contest* and *dissolves* in a genuinely multi-round one. The founding qualifier of this project, that a toehold’s value compounds across three or more rounds, had the right axis and the wrong sign: the effect does not compound, it evaporates. One caveat must be stated: at a *zero* toehold the two-move game already carries two equilibria, so multiplicity is not created by adding rounds in the abstract; it is created by adding rounds *at a positive toehold*, which is the case of interest.

7 Reconciliation and the toehold puzzle

These findings do not contradict the literature; they sharpen it. For Fishman [1988], preemptive jump bidding is a *signalling* phenomenon, not a toehold phenomenon, and the mechanism we find is the same winner’s-curse inference: the toehold-holder jump-bids, and a rival holding a weak signal infers that competing means paying at least the jump and winning mainly where its own signal is wrong, so it folds. No toehold is required for that, which is exactly why preemption appears at $\theta = 0$. What dies is only the narrower, Bulow-Huang-Klemperer-flavored claim that a *toehold* adds a deterrence channel on top of the signalling one. We therefore reproduce Fishman-style signalling preemption computationally and show the toehold does not add to it.

This speaks to the toehold puzzle. If a toehold buys profit but no identified deterrence, then part of the rarity that Betton et al. [2009] document may need no appeal to hidden costs: in our model the theoretical bidding edge that toeholds are supposed to confer is not identified once the contest runs past one round. This is a model-side reason the edge may be weaker than assumed, complementary to the cost-based explanations rather than a replacement for them, and it does not speak to the selection correction of Dai et al. [2021], which is an empirical matter our computation cannot settle. The aggressive-toehold equilibrium requires the rival to *read* a

jump bid as strength; absent that coordination, the toehold still raises the holder’s profit but buys no deterrence.

What may and may not be claimed. We can say that preemption is *sustainable* in the sequential contest and *absent* in the sealed auction of our companion paper (whose equilibrium bid is flat in the toehold on the identical grid). What that demonstrates is the structural half of the companion paper’s assertion: a sequential contest can express preemption where a sealed one cannot. It does *not* vindicate the *toehold* half of that assertion, which our zero-toehold result contradicts. and that the toehold-holder’s equilibrium profit rises robustly with the toehold. We do *not* claim that a toehold implies any particular amount of deterrence, or any smooth deterrence-versus-toehold curve. The deterrence prediction is not identified, and drawing that curve from a single family of solves would be precisely the plausible-but-wrong result our methodology is built to avoid.

8 Benchmarking solvers on the sequential game

We next ask which solvers actually solve the sequential game, on commodity CPU with no frontier-model spend. On a three-round instance small enough that exact counterfactual regret minimization and exact exploitability remain a valid reference (num_values = 3, num_bids = 6, about 6,655 states), we run deep self-play for PPO, PPG, and deep policy gradient, alongside NFSP, and compare against CFR, PSRO, and Deep CFR, scoring each by exact exploitability of its tabular tail-average over five seeds (Figure 2).




Figure 2: Deep self-play on the three-round takeover auction (mean $\pm$ std over five seeds). PPG converges to a near-exact equilibrium (0.0007 exploitability), below the exact-CFR reference at the budget shown; PPO is competitive but seed-sensitive; NFSP does not converge within the shared budget.

The ordering is stark. PPG reaches 0.0007 ± 0.0002 exploitability, with all five seeds in agreement, dipping below the exact-CFR reference of 0.0069 at the budgets shown (we do not read this as PPG beating CFR at equilibrium-finding: CFR run longer would go lower, and it reached its value faster; rather, PPG converges to a genuine near-exact equilibrium and is the strongest deep method by a wide margin). PPO is competitive at 0.031 but seed-sensitive (± 0.029). Deep policy gradient plateaus near 0.37. NFSP, a standard imperfect-information

baseline, does not converge within the shared 300,000-episode budget, crawling from 1.83 only to 1.16; it was already the weakest method on our sealed game, and the sequential structure amplifies the gap. Absolute exploitability is not comparable across the sealed and sequential games; only the within-game ordering is.

Scaling in the number of rounds. Because the ascending-bid ceiling caps tree depth at the price grid, we widen the grid (`num_bids` = 8) so that added rounds genuinely deepen the tree, and measure the wall-clock each solver needs to reach an exploitability of 0.05 as the contest grows from one to four rounds (Figure 3). The tree grows about eighteen-fold over the range. Exact CFR remains faster than PPO at every round count in this tractable range, but its cost climbs steeply from one round to two and then plateaus (5.4, 30.2, 37.7, 34.0 seconds) while PPO’s stays roughly flat (103, 137, 136, 124 seconds), so the ratio of CFR to PPO wall-clock narrows from 0.05 at one round to 0.27 at four: the gap closes but does not cross. The crossover belongs to the regime where the game can no longer be enumerated at all.




Figure 3: Scaling in the number of rounds. Left: game-tree size grows steeply with rounds (the flattening at $R = 4$ is the price-grid ceiling beginning to bind). Right: exact CFR’s wall-clock cost to reach exploitability 0.05 jumps from one round to two and then plateaus, while PPO’s is roughly depth-invariant.

9 The intractable regime

Both the number of rounds and the number of private signals enlarge the game, but only the signals do so without bound: information sets grow as num_values^k in the signal count k , and the history count as num_values^{1+2k} times the bidding subtree, so a multi-round, multi-signal instance quickly outruns exact enumeration. On an instance with $k = 6$ signals per bidder over three rounds (on the order of 10^7 histories, beyond exact CFR and exact exploitability), we train PPO and PPG self-play without tabulating and score them with a learned-best-response exploitability estimate: freeze the policy, train a PPO best-responder against it, and Monte-Carlo the gain from deviating. A learned best response underestimates the true one, so the estimate is a lower bound, which we report as such rather than as a Nash certificate. Against this estimator, uniform-random play is highly exploitable (1.77) and a naive no-shading bidder economically so (0.083), while PPO and PPG drive the estimate to 0.004 ± 0.003 and 0.003 ± 0.004 over three seeds, below the naive bidder and down to the estimator’s calibrated resolution

floor, consistent with near-unexploitable play, in a regime where the exact solvers cannot run (Figure 4). The two learners are not separable from each other at that spread; both are simply at the floor.




Figure 4: Intractable multi-round, multi-signal auction ($k = 6$ signals, three rounds, about 10^7 histories). PPO (0.004 ± 0.003) and PPG (0.003 ± 0.004) drive the learned-best-response exploitability estimate below a naive unshaded bidder (0.083) and far below uniform play (1.77), in a regime where exact solvers cannot run.




Figure 5: Estimator calibration on enumerable multi-signal sequential games ($k = 1, 2$). The learned-best-response estimate tracks exact exploitability across policies of known exploitability (CFR equilibrium mixed with uniform play), lying close to the $y = x$ line over the whole range.

To show the estimate is trustworthy on the sequential multi-signal information-set structure, and not merely reading a weak best response, we calibrate it on enumerable instances ($k = 1, 2$) where exact exploitability is available as ground truth: we build policies of known exploitability by mixing the CFR equilibrium with uniform play and confirm the learned estimate tracks the exact value (Figure 5). The tracking is tight across the whole range: at $k = 2$, for example, the estimate reads 0.257 against an exact 0.270 at a 40% uniform mix and 0.765 against an exact 0.764 at full uniform, and at the CFR equilibrium both fall to the estimator’s resolution floor of a few thousandths. The estimate is therefore a lower bound calibrated on the sequential multi-signal structure at $k = 1, 2$, not an artifact of a weak best-responder. It is applied at $k = 6$, where no exact ground truth exists to calibrate against; the calibration is evidence that the estimator tracks the truth on this game’s information-set structure, not a guarantee at a signal count we cannot enumerate.

10 Discussion and limitations

The economic headline is a cautionary one for practice and theory alike: the deterrence value of a toehold, the thing that would justify acquiring one, is a feature of our model truncated at a single round and does not survive a genuine second round; by the third the equilibria multiply and it is not identified at all. This does not say toeholds are worthless (they still raise the holder’s profit), but it removes the identified strategic edge that the puzzle literature has struggled to reconcile with their rarity.

Several limitations bound the claims, and they are worth being blunt about.

One instance. Every economic result above is one parameterization: three value levels, a nine-level price grid, one signal per bidder, and symmetric signal noise of 0.5. We sweep the toehold, the number of rounds, and the grid resolution, and nothing else. We have not varied the

number of values, the number of signals, or the noise, so we do not know whether the multiplicity or the two-move saturation survives a richer signal space. Where we write that the deterrence channel is an artifact of stopping at one or two moves, that is a statement about *our* model truncated at one or two moves, not a theorem about the class of short-contest models; we solve neither Dodonova’s model nor Bulow-Huang-Klemperer’s, and reproducing their comparative static at $R = 1$ is not the same as deriving their result.

Per-configuration, not theorems. The value-is-pinned observation is verified where measured and explicitly not a theorem: it fails at one and two rounds with a zero toehold, where multiple values coexist. Nor is the non-identification an econometric claim. What we exhibit is a set of computed equilibria at one instance that a single deterrence number cannot summarize; that is model-side non-identification, and it does not by itself establish that deterrence is unidentifiable in field data.

Precision. The ε -equilibria we exhibit are certified to ε between 2×10^{-6} and 8×10^{-5} , small but not exact; in a flat game an ε -equilibrium can sit some distance from an exact one, so we phrase every claim in terms of ε -equilibria that a referee can reproduce. The intractable-regime numbers are learned-best-response *lower bounds*, calibrated at $k = 1, 2$ and applied at $k = 6$, not certificates.

Grid robustness is not established. We had hoped to report that the multiplicity survives grid refinement. It does not survive the test, in the weaker sense that the test does not run: in the released refinement study only the base nine-level grid reaches the acceptance bound for most restarts. At seven levels five of eighteen solves are accepted and none of them exhibits the low-deterrence conduct; at thirteen levels one of eighteen is accepted. So the multiplicity is exhibited at one grid resolution, and whether it is a discretization artifact is neither established nor refuted. The finer cells are the slowest solves in the suite, at twelve to twenty hours each, and are left to future compute. The headline does not rest on this study, which rests instead on the rounds sweep at $\varepsilon \sim 10^{-6}$, but the gap is real and we do not want it read as robustness.

The mechanism figure lags the tables. The rival policies behind Figure 1 in the released artifacts are default-budget solves at NashConv 10^{-4} , not probe-precision ones, and at two toeholds the least-detering equilibrium found is not a passive rival at all. The forced-opening construction itself is exact and the control panel is solver-free, so the mechanism reading is sound where a passive rival exists; the certified re-run of that study is wired into **certify** but its output is not what the committed figure shows.

11 Conclusion

We solved a multi-round ascending takeover auction with a toehold and found that the toehold’s supposed deterrence edge is not what three decades of two-move intuition suggests. Equilibria are multiple and give the same profit with different conduct; preemption occurs even without a toehold; and the clean toehold-to-deterrence relationship is an artifact of stopping at two moves, dissolving as soon as the contest has a genuine second round and leaving deterrence unidentified by the third. Preemptive jump bidding is real and computationally reproducible, but as Fishman-style signalling, not as a toehold effect, which rationalizes the toehold puzzle. Methodologically, forced-opening best responses read the economics without trusting any single solver, and a fictitious-play solve taken at face value manufactures confident artifacts. On the computational side, a generic policy-gradient method solves the sequential game where a standard imperfect-information baseline does not, and the analysis extends into a regime beyond exact enumeration. We release the games, solvers, and experiments so the results can be reproduced and extended.

References

- Christopher Avery. Strategic jump bidding in English auctions. *Review of Economic Studies*, 65(2):185–210, 1998.
- Sandra Betton, B. Espen Eckbo, and Karin S. Thorburn. Merger negotiations and the toehold puzzle. *Journal of Financial Economics*, 91(2):158–178, 2009. doi: 10.1016/j.jfineco.2008.02.004.
- Noam Brown, Adam Lerer, Sam Gross, and Tuomas Sandholm. Deep counterfactual regret minimization. In *International Conference on Machine Learning (ICML)*, 2019. arXiv:1811.00164.
- Jeremy Bulow, Ming Huang, and Paul Klemperer. Toeholds and takeovers. *Journal of Political Economy*, 107(3):427–454, 1999.
- E. C. Capen, R. V. Clapp, and W. M. Campbell. Competitive bidding in high-risk situations. *Journal of Petroleum Technology*, 23(6):641–653, 1971.
- Karl Cobbe, Jacob Hilton, Oleg Klimov, and John Schulman. Phasic policy gradient. In *International Conference on Machine Learning (ICML)*, 2021. arXiv:2009.04416.
- Yun Dai, Sebastian Gryglewicz, and Han T. J. Smit. Less popular but more effective toeholds in corporate takeovers. *Journal of Financial and Quantitative Analysis*, 56(1):283–312, 2021. doi: 10.1017/S0022109019001029.
- Kent D. Daniel and David Hirshleifer. A theory of costly sequential bidding. *Review of Finance*, 22(5):1631–1665, 2018.
- Greg d’Eon, Neil Newman, and Kevin Leyton-Brown. Understanding iterative combinatorial auction designs via multi-agent reinforcement learning. In *Proceedings of the 25th ACM Conference on Economics and Computation (EC ’24)*, 2024. doi: 10.1145/3670865.3673644.
- Anna Dodonova. Toeholds and signalling in takeover auctions. *Economics Letters*, 117(2):386–388, 2012. doi: 10.1016/j.econlet.2012.06.023.
- Michael J. Fishman. A theory of preemptive takeover bidding. *RAND Journal of Economics*, 19(1):88–101, 1988.
- Johannes Heinrich and David Silver. Deep reinforcement learning from self-play in imperfect-information games. In *NIPS 2016 Deep Reinforcement Learning Workshop*, 2016. arXiv:1603.01121.
- Paul Klemperer. Auction theory: A guide to the literature. *Journal of Economic Surveys*, 13(3):227–286, 1999.
- Marc Lanctot, Vinicius Zambaldi, Audrunas Gruslys, Angeliki Lazaridou, Karl Tuyls, Julien Pérolat, David Silver, and Thore Graepel. A unified game-theoretic approach to multiagent reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017. arXiv:1711.00832.
- Marc Lanctot, Edward Lockhart, Jean-Baptiste Lespiau, et al. Openspiel: A framework for reinforcement learning in games. *arXiv preprint arXiv:1908.09453*, 2019.
- Paul R. Milgrom and Robert J. Weber. A theory of auctions and competitive bidding. *Econometrica*, 50(5):1089–1122, 1982.
- Zain Naboulsi. How much due diligence before you bid? learning in intractable takeover auctions, 2026. arXiv preprint; identifier to be added upon posting.

- Max Rudolph, Nathan Lichtlé, Sobhan Mohammadpour, Alexandre Bayen, J. Zico Kolter, Amy Zhang, Gabriele Farina, Eugene Vinitzky, and Samuel Sokota. Reevaluating policy gradient methods for imperfect-information games. In *International Conference on Learning Representations (ICLR)*, 2026. arXiv:2502.08938.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Samuel Sokota, Ryan D’Orazio, J. Zico Kolter, Nicolas Loizou, Marc Lanctot, Ioannis Mitliagkas, Noam Brown, and Christian Kroer. A unified approach to reinforcement learning, quantal response equilibria, and two-player zero-sum games. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2206.05825.
- Robert Wilson. A bidding model of perfect competition. *The Review of Economic Studies*, 44 (3):511–518, 1977.
- Martin Zinkevich, Michael Johanson, Michael Bowling, and Carmelo Piccione. Regret minimization in games with incomplete information. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2007.